

BoardroomAI — session architecture

The model is called once per turn. Turn-taking, quorum and vote tallying are deterministic TypeScript.

BROWSER

Boardroom UI

Drives the turn loop · speaks each turn · shows the scoring that chose the speaker

boardroom-session.tsx

NEXT.JS ROUTE HANDLERS · THE ONLY PLACE SECRETS ARE READ

POST /api/pitches

Creates the meeting

POST .../debate

One executive turn

POST .../finalize

Votes · report

POST /api/speech

Edge neural TTS

ORCHESTRATION — DETERMINISTIC TYPESCRIPT · ~2,300 LINES · NO MODEL CALLS

topics.ts

137 weighted terms · 6 topics
Founder words weigh 2.5x

expertise.ts

Topic authority per seat
Non-zero floor everywhere

meeting-memory.ts

Positions · open challenges
Derived, never stored

debate-policy.ts — who speaks next

0.45 relevance + 0.20 fairness + 0.15 mention + 0.10 challenge + 0.10 priority
3 phases · shared byte-for-byte with the browser

fact-check.ts

Extracts numeric claims,
matches them to sources

MODELS & RETRIEVAL · INFERENCE ONLY, NO TRAINING

llama-3.1-8b-instant

~16 debate turns a session

llama-3.3-70b-versatile

2 calls: verdict + analysis

Tavily search

Research + VC seats only

PERSISTENCE

Supabase Postgres

meetings · messages (+ verification) · votes (+ 6 decision fields) · reports · research_findings
Row-level security on every table — the database itself refuses another user's rows

turn loop · x16

— accent = our orchestration and our loop, not the model's

The model produces one executive's dialogue for one turn. Everything else is code.